

Golden Thread — Getting Started

Reader: the first thirty minutes — the happy path only **Claims last checked against the code:** 2026-10-01 (gt 0.18.1) — see *The documents, and what belongs in each* in [CLAUDE.md](#) .

A guided walkthrough for your first session. Six steps, ~15 minutes.

Requirements: Python 3.8+, Claude Code installed. Written against **gt v0.18.1**.

Step 1 — Install

Unzip the plugin and run the installer:

```
unzip golden-thread-plugin.zip
cd golden-thread-plugin
bash install.sh --vault ~/Documents/GoldenThread
```

Then **restart Claude Code**. The `/gt:` skills are not available until you do.

`--vault` names the folder your memory will live in, and lets the installer finish the job: it creates the vault (or connects one you already have) and wires the Core-rule enforcement hooks, which are wired against a vault and cannot be wired without one.

Leave it off and the installer asks, if you are at a terminal. If something else is running it — an agent, CI — it stops with exit 4 and asks for the decision rather than choosing a directory for you. `bash install.sh --no-vault` installs the plugin alone, deliberately.

The installer sets up **gt** — memory, projects, session workflow, and enforcement hooks — plus each optional **module** that is on: - **gt-wiki** — knowledge base with immutable sources - **gt-demo** — the guided tour, in its own throwaway vault - **gt-watch** — follow git repos you depend on (`/gt-watch:gt-watch`) - **gt-report-card** — the session report card at `/compact` - **gt-flow** — a picture of how your knowledge moved (`/gt-flow:gt-flow`) - **gt-visualize** — a 3D walkthrough of how a codebase works, or a code city of it (`/gt-visualize:gt-visualize`) - **gt-farm** — work packets for an external AI; **off** unless you add it with `bash install.sh --with farm` (or had `/gt:gt-farm` before upgrading)

`bash install.sh --list-modules` shows each one and why it is on or off; `--without <name>` leaves one out, and the choice is remembered.

Step 2 — Initialize your vault

Open any Claude Code session and run:

```
/gt:gt-init
```

Claude will ask for: 1. A vault path — an Obsidian vault folder is ideal, but any folder works 2. Your domain/team name — e.g. "Platform Engineering" or "My Team"

The script creates the folder structure, writes `~/ .claude/vault-config.json` (the pointer every other skill reads), and installs **Core-rule enforcement hooks** into `~/ .claude/settings.json`. Safe to re-run on a machine you've already set up.

If you already have an Obsidian vault, just give `gt-init` that path. It will only add what's missing — your existing notes are untouched.

Looking at the vault in Obsidian (optional). A vault is a folder, so there is no import: install Obsidian, choose **Open folder as vault**, and pick this folder. The vault's own `OPEN-IN-OBSIDIAN.md` has the path and the details. Install **Dataview** if you want to sort tasks by the `[p:: 1]` inline fields, and **Obsidian Git** if you want your hand edits committed on a timer alongside Claude's.

Step 3 — Create your first project

```
/gt:gt-create
```

Claude will ask for a project name and a quick brain dump of what it is. That brain dump becomes `idea.md` — the immutable origin story for the project.

For a project you're already working on, use:

```
/gt:gt-ingest
```

This scans your existing notes and memory files and imports them into the vault without touching the originals.

Step 4 — Open a project at the start of each session

```
/gt:gt-open <project-name>
```

This reads your project docs in order — `idea` → `research` → `decisions` → `design` → `spec` → `runbook` — reads `CONVENTIONS.md` and `PROTOCOL.md` once per session, then reads the session memory index and stops. It summarizes where things stand and asks where to pick up.

Why it stops at the index: a project with 70 notes would otherwise dump 2,000 lines into context. `gt-open` reads the one-line index and loads individual files only when you need them. If `research.md` exceeds ~200 lines, it reads its one-line-per-entry digest (`research-digest.md` , kept current by `/gt:gt-work`) or, failing that, the headings, then only the relevant sections. Come back after a week away and it opens with a short generated catch-up first.

Step 5 — Write back at the end of each session

```
/gt:gt-work
```

This is the step that keeps the vault alive. Claude will write back what was learned:

DESTINATION	WHAT GOES THERE
<code>research.md</code>	New findings and gotchas, dated
<code>decisions.md</code>	Architectural choices (ADRs)
<code>design.md</code>	Updated architecture if it changed
<code>runbook.md</code>	Operational commands and landmines
<code>memory/*.md</code>	Session state

`gt-work` also checks whether any findings should be promoted wider (to `Knowledge/` or `global-memory/`) and asks you before writing. It updates the project's `README.md` frontmatter properties (`stage` , `topology`) if the project changed phase.

Skipping `gt-work` is how the vault decays. Facts that aren't written down are forgotten by the next session.

Step 6 — Capture from anywhere, then ask what's next

Steps 4 and 5 keep one project honest. This step is what keeps *all* of them in view, and it is the step that matters most if your attention is split across many things.

Capture wherever you are. You are working in project A and a thought about project B arrives. Do not switch projects. Add one checkbox line to `<vault>/INBOX.md` :

```
- [ ] Chrome extension: the confirm dialog retry may be resubmitting instead of confirming
```

No project, no priority, no date. The rollup shows every unchecked inbox line at the top of `TASKS.md`, so it cannot be lost, and nowhere else, so it does not rank. (If you keep Obsidian daily notes, the sweep reads those too.)

Sweep the inbox. Once a week, or whenever the notes pile up:

```
/gt:gt-review
```

Claude reads the inbox, asks where each line belongs, and files it: a task under an existing project, a sub-project, or a new project with its own `idea.md`. The inbox line is checked off with a `→ [[slug]]` pointer so the same thought is never surfaced twice. An idea becomes a `p:: 3` task with no due date.

Ask what's next. Every project keeps its open tasks under `## Tasks` in its `README.md`, one checkbox per task, with optional fields:

```
- [ ] Reproduce the dialog retry against a live page [p:: 1] [waiting:: agent] [since:: 202
```

One command rolls every project's tasks into a single ranked list:

```
python3 <vault>/Projects/golden-thread/tools/gt_tasks.py # then read TASKS.md
```

Rank is computed against the clock, not stored: a `p:: 1` older than seven days escalates, a `due` inside three days escalates. **Regenerate before reading** — a stale `TASKS.md` is last week's ranking. `waiting:: user` is the list of things blocked on you; `waiting:: agent` is what the next session should pick up.

That is the loop: capture without switching, sweep to file, regenerate to decide.

Close what is finished. The rollup's *Review* section, the `/compact` report card (report-card module) and `gt-work` all ask the same question when a project's signals say it may be done: most tasks past due, most tasks checked, three quiet weeks, or nothing open. Answer yes, no or later; every answer is recorded so the question learns your pattern. Closing sets `stage: complete` and shelves leftover tasks at `p:: 7` — they stay in the README, they just stop outranking live work.

Typical session pattern

```
/gt:gt-open my-project    ← load context, see where you left off

# ... work happens ...

/gt:gt-route              ← lost? drifted? "where is this going, and where
                           does it belong?" Safe to run any number of times

/gt:gt-work                ← write back findings before closing
```

Periodic maintenance

```
/gt:gt-lint                ← find broken links, orphaned pages, unlisted memory files, unen
/gt:gt-promote             ← graduate a finding that proved true across projects
/gt:gt-refresh             ← check if any source documents changed upstream
/gt:gt-review              ← file the inbox (INBOX.md, plus daily notes if you keep them)
/gt:gt-settings            ← view and toggle what Golden Thread does automatically
/gt:gt-validate            ← independently verify a claim before recording it as fact

# from modules
/gt-watch:gt-watch         ← watch a git repo you depend on; a security fix upstream opens
/gt-flow:gt-flow           ← draw how knowledge climbed the ladder, one offline HTML file (
/gt-visualize:gt-visualize ← explain a codebase in 3D, or draw it as a code city (one offli
/gt-farm:gt-farm           ← route bulk or mechanical tasks to an external AI service as a
/gt-demo:gt-demo           ← guided tour in a throwaway demo vault, nine acts plus one per
```

Adding a knowledge base (gt-wiki)

If you want a structured wiki alongside your project memory:

```
/gt-wiki:gt-wiki-init
```

Then add sources:

```
/gt-wiki:gt-wiki-ingest <url or file path>
```

Query it:

```
/gt-wiki:gt-wiki
```

Check its sources for upstream changes:

```
/gt-wiki:gt-wiki-refresh
```

The wiki uses immutable sources — every ingested document is stored verbatim and never modified. Knowledge pages are synthesized summaries that link to the originals. When a source changes upstream, refresh adds a new source that supersedes the old one; for sources in a local git repo the change is detected by `git diff`, not by eye.

Key concepts in one paragraph

The vault is organized by **how widely a fact applies**. Session findings go in `memory/`. Settled discoveries go in `research.md`. Choices that shaped the architecture go in `decisions.md`. Facts true across multiple projects go in `Knowledge/`. Facts needed in every session everywhere go in `global-memory/`. Above all of these sits the **Core-rule tier**: rules enforced via hooks wired into `~/.claude/settings.json`, not merely stored as files. `gt-init` installs these hooks automatically. Promote a fact when a second project proves it general — not before. When unsure, file it narrow; promotion is cheap, demotion isn't. One more principle: **claims should be verified by re-deriving them, not by reviewing the reasoning that produced them** — a reviewer who sees your reasoning inherits your blind spot. Use `/gt:gt-validate` before recording any finding that matters.

Troubleshooting

Skills not showing up — restart Claude Code after install.

"vault-config.json not found" — run `/gt:gt-init`.

A rule seems ignored — `CLAUDE.md` is read at session start; rules added mid-session apply from the next session.

Memory files aren't loading — that's by design. Ask for a file by name, or run `/gt:gt-lint` to check if it's unlisted.

Core rules seem to apply sometimes but not always — run `/gt:gt-lint` and look for `core-unenforced` findings. The hook may not be wired. Fix:

```
python3 <plugin>/scripts/vault_init.py install-core-rules --vault <vault>
```

I want to turn off the automatic session report — run `/gt:gt-settings set report_card off`, or remove the module with `bash install.sh --without report-card`.

Timestamp missing from responses — the enforcement hook is unwired. Run `/hooks` in Claude Code or restart. If it persists, re-run `install.sh`.

For the complete skill reference, see [MANUAL.md](#).